

Cluster Computing: SMP Assignment

Martin Rehr

eScience Center
Niels Bohr Institute
University of Copenhagen

Copenhagen, February the 25th, 2013

CT-Reconstruction

Using Symmetric Multi Processors (SMP)

Outline

- ➊ Computed Tomography (CT)
- ➋ CT reconstruction
- ➌ SMP Assignment
- ➍ Programming SMPs

CT Scanning

We go to the hospital, right ?!?

CT Scanning: Veterinary



CT Scanning: Danish Crown



CT Scanning: NOVO



CT Scanning: Grundfos



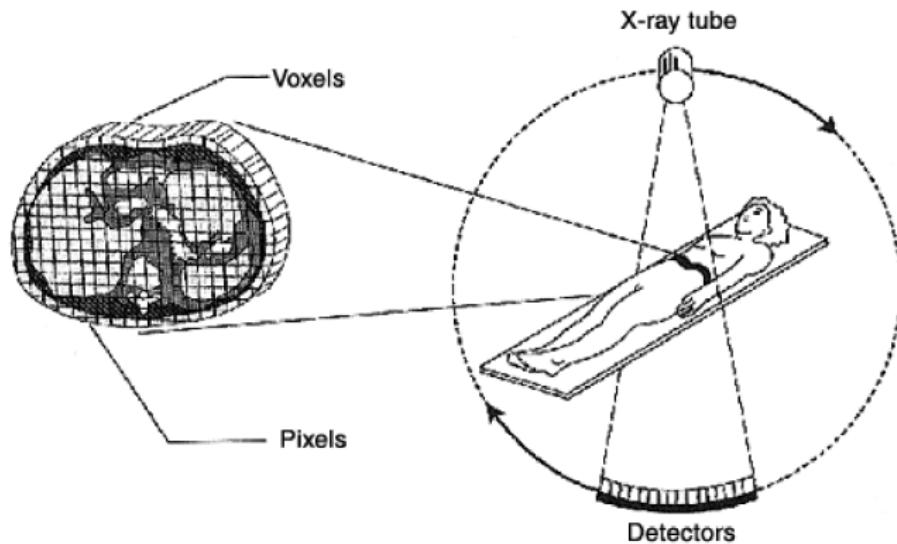
CT Scanning: NBI XRay



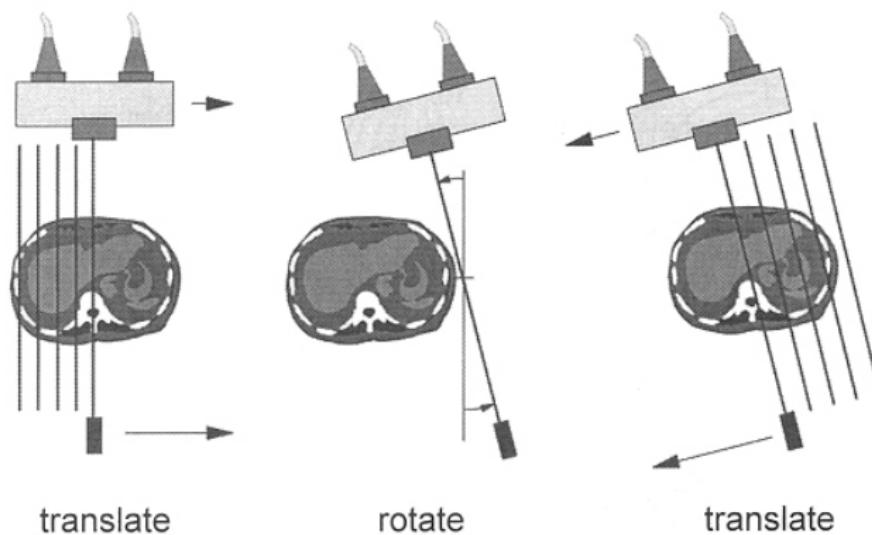
John Schinnerer www.IllustrationOnline.com

Computed Tomography

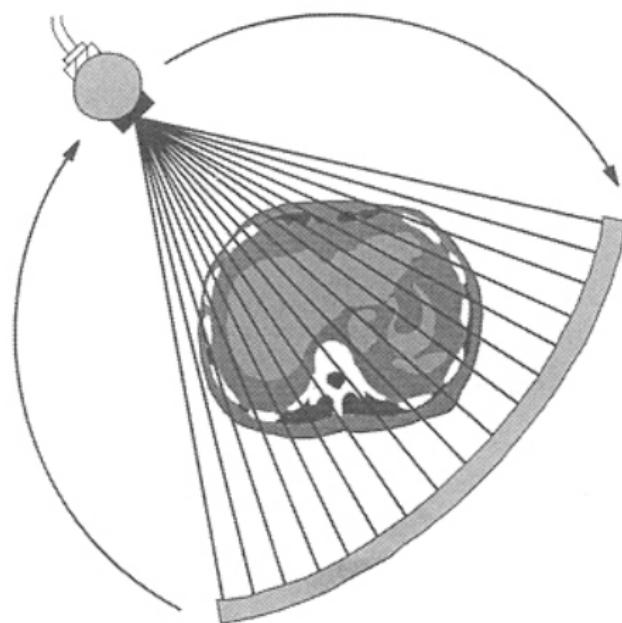
Overview



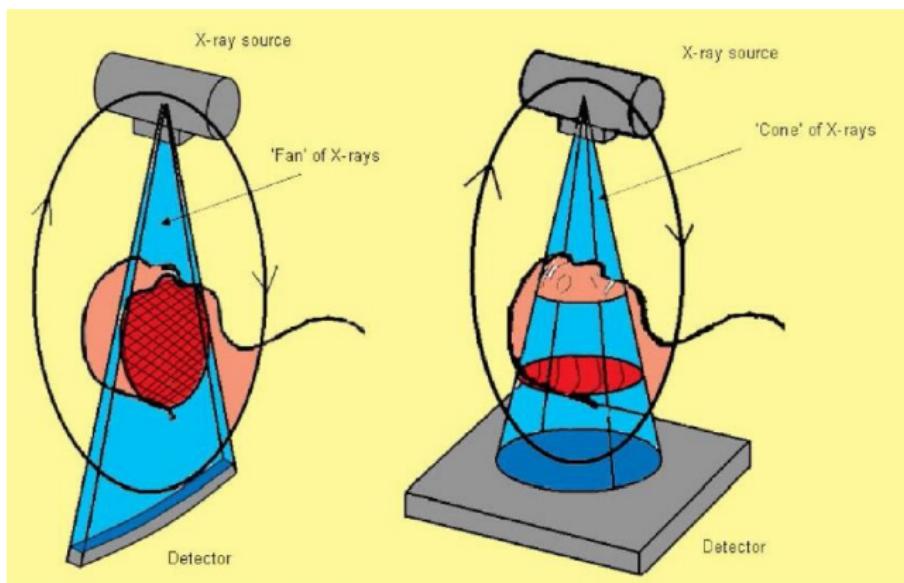
Parallel Beam



Fan Beam

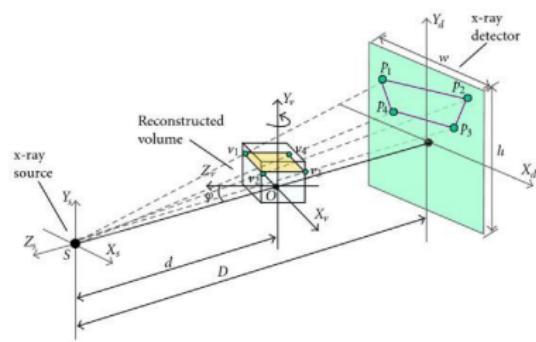
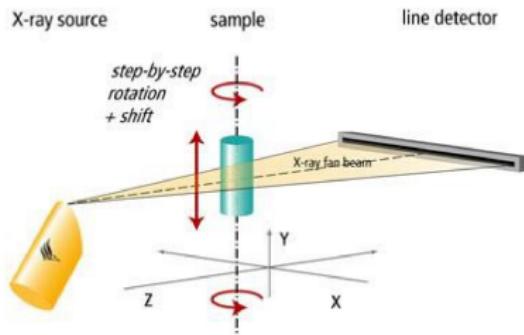


Fan vs. Cone Beam

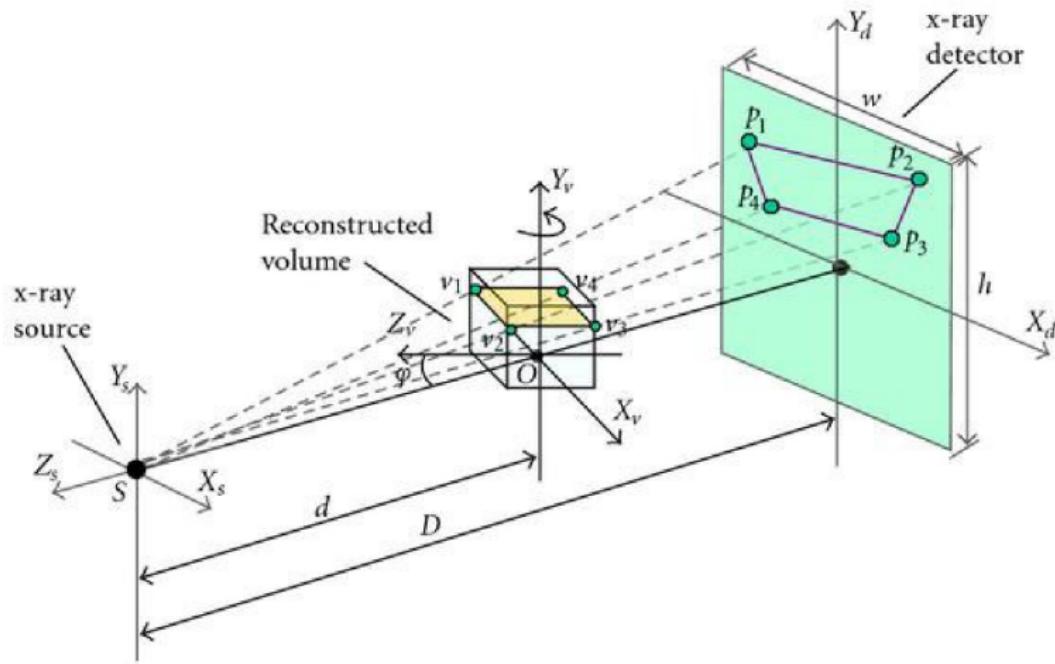


© J Can Dent Assoc 2006; 72(1); 75-80

Fan vs. Cone Beam

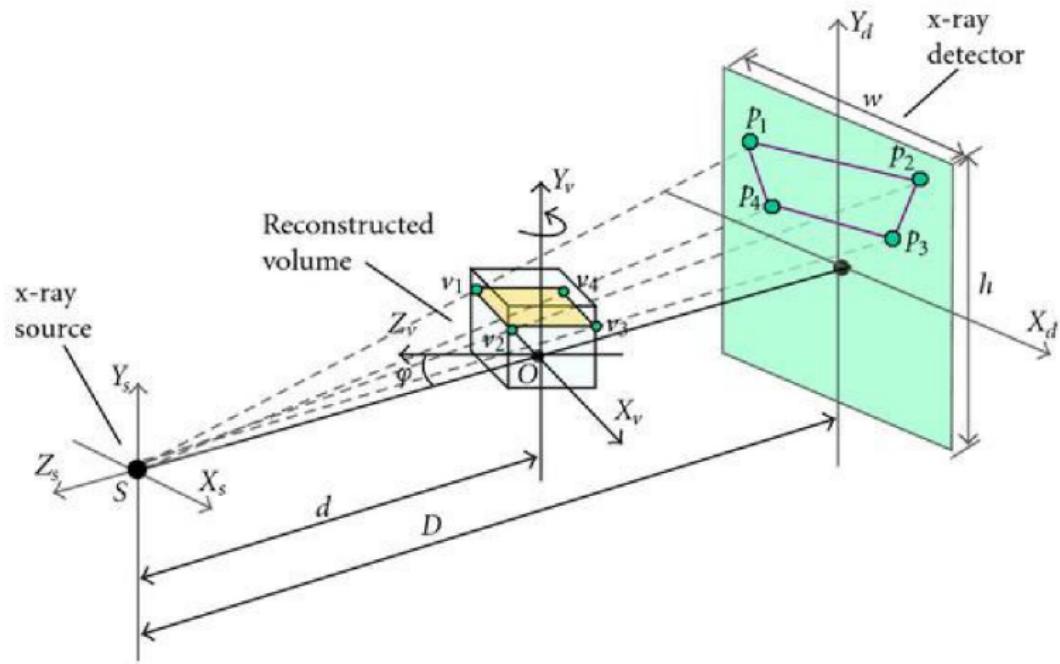


Cone Beam

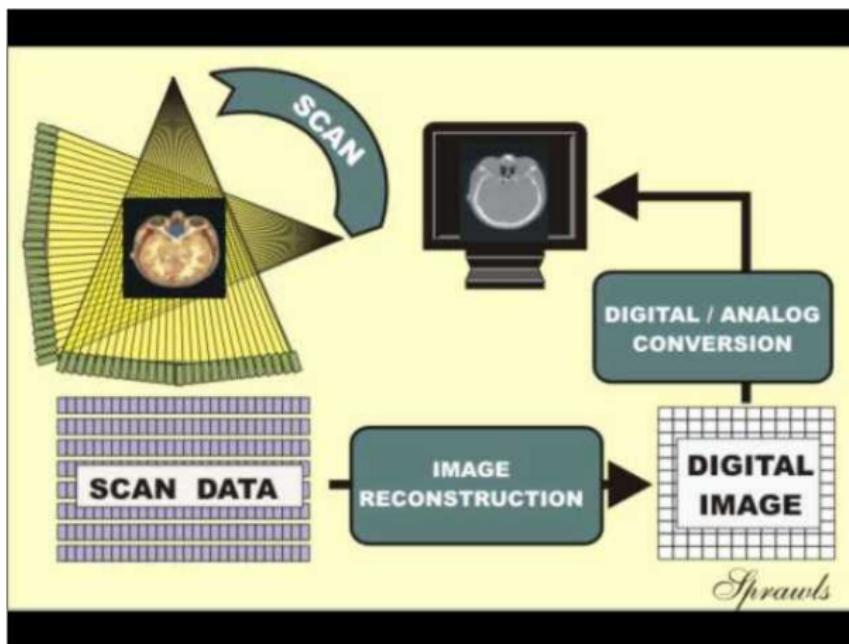


CT Reconstruction

CT Reconstruction



CT Reconstruction



CT Reconstruction

skull-cleaned_projections_...
69/320; 256×192 pixels; 32-bit; 60MB

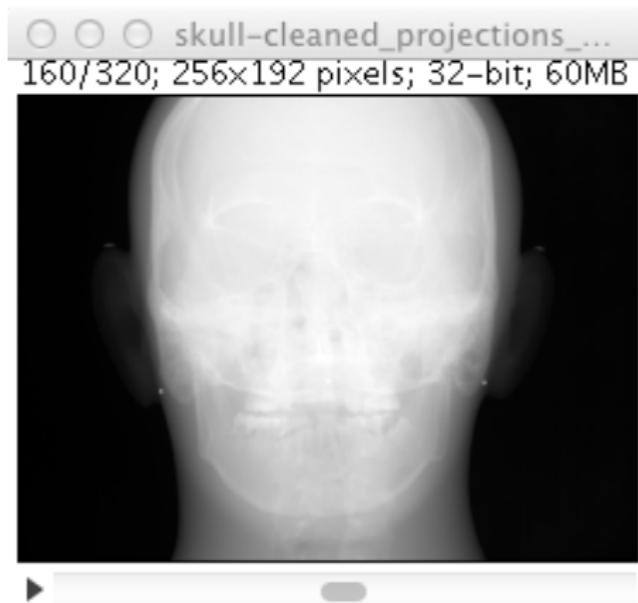


CT Reconstruction

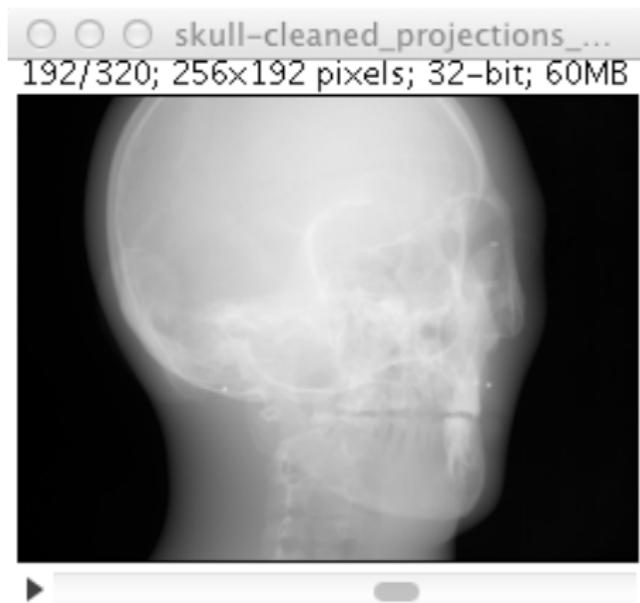
skull-cleaned_projections_...
118/320; 256x192 pixels; 32-bit; 60MB



CT Reconstruction

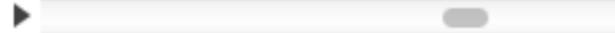


CT Reconstruction

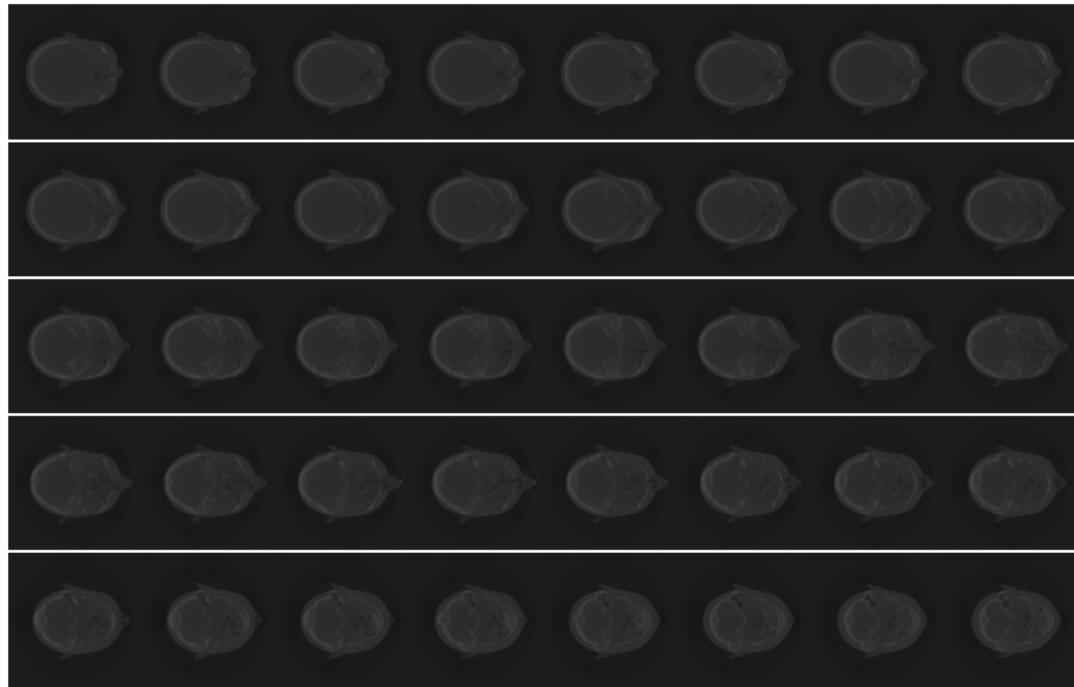


CT Reconstruction

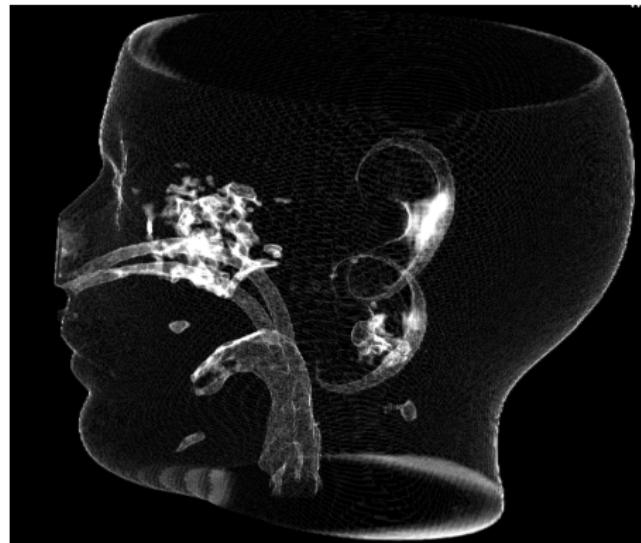
skull-cleaned_projections_...
241/320; 256x192 pixels; 32-bit; 60MB



CT Reconstruction



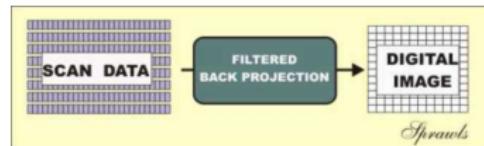
CT Reconstruction



Feldkamp, Davis, and Kress (FDK)

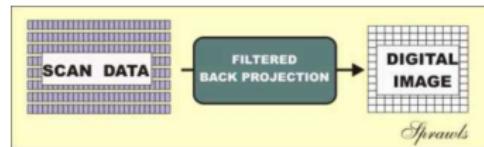
Cone Beam Reconstruction

- Projection Weighting
- Projection Filtering
- Back Projection
- Volume Weighting



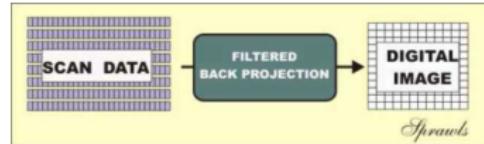
Cone Beam Reconstruction

- Projection Weighting
- Projection Filtering
- Back Projection
- Volume Weighting



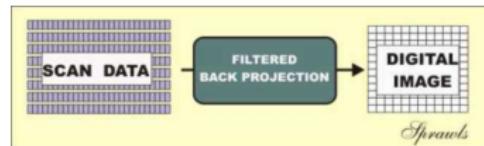
Cone Beam Reconstruction

- Projection Weighting
- Projection Filtering
- Back Projection
- Volume Weighting



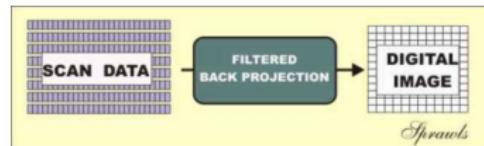
Cone Beam Reconstruction

- Projection Weighting
- Projection Filtering
- Back Projection
- Volume Weighting

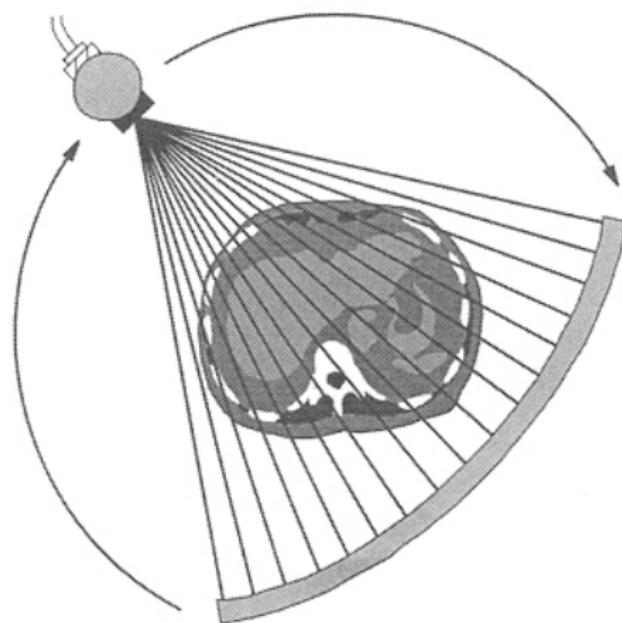


Cone Beam Reconstruction

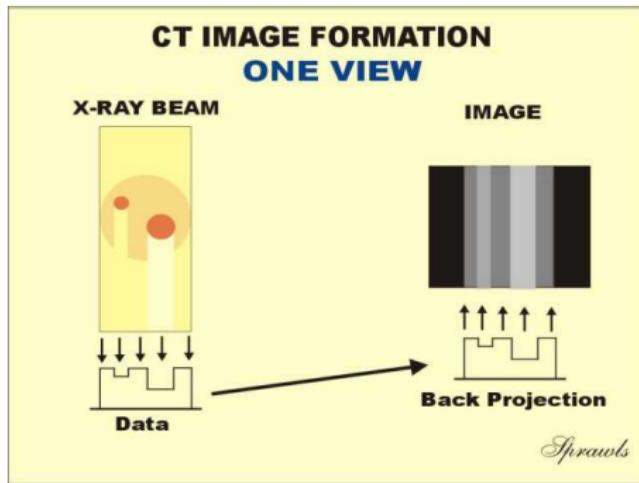
- Projection Weighting
- Projection Filtering
- Back Projection
- Volume Weighting



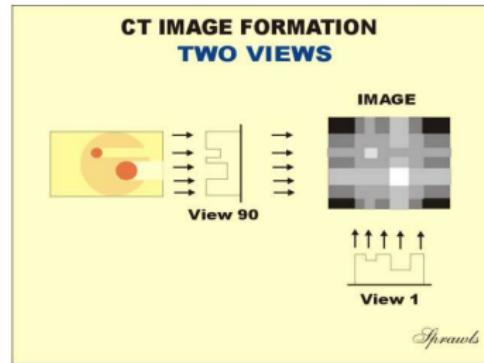
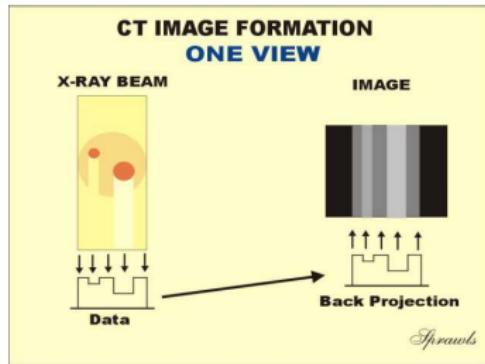
Projection Weighting



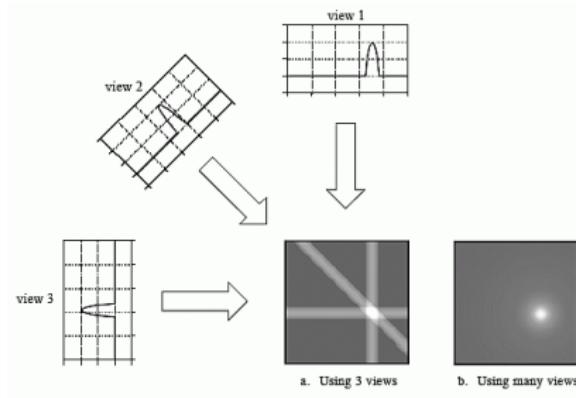
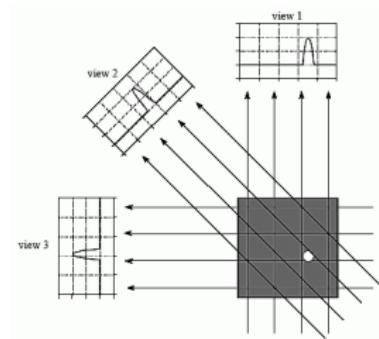
Back Projection



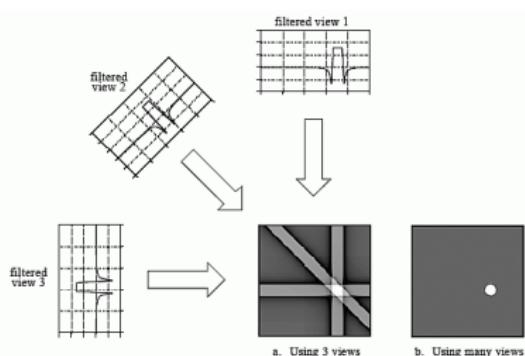
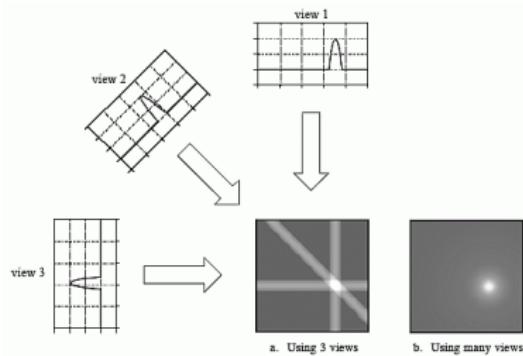
Back Projection



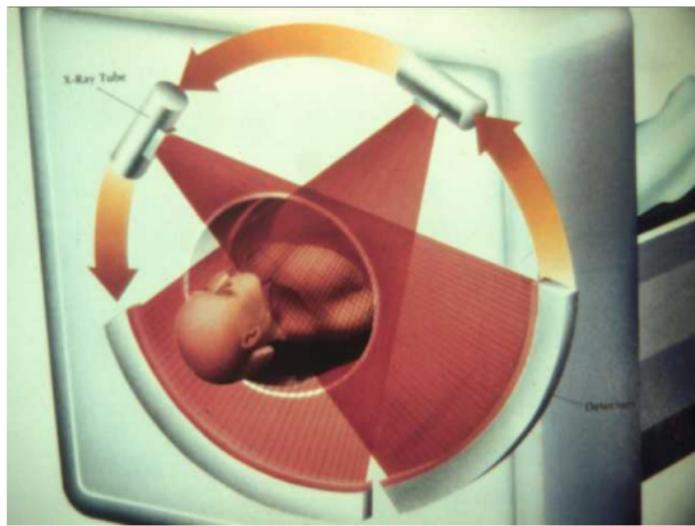
Back Projection



Filtered Back Projection



Volume Weighting



Assignment

Make a SMP version of FDK

Assignment files

- fdk.py
 - Main file used for benchmarks
- fdkvisual.py
 - Main file used for visualization
- fdkload.py
 - Module used for loading data
- fdkcore.py
 - Core reconstruction kernel to be modified

Assignment files

- fdk.py
 - Main file used for benchmarks
- fdkvisual.py
 - Main file used for visualization
- fdkload.py
 - Module used for loading data
- fdkcore.py
 - Core reconstruction kernel to be modified

Assignment files

- fdk.py
 - Main file used for benchmarks
- fdkvisual.py
 - Main file used for visualization
- fdkload.py
 - Module used for loading data
- fdkcore.py
 - Core reconstruction kernel to be modified

Assignment files

- fdk.py
 - Main file used for benchmarks
- fdkvisual.py
 - Main file used for visualization
- fdkload.py
 - Module used for loading data
- fdkcore.py
 - Core reconstruction kernel to be modified

Assignment files

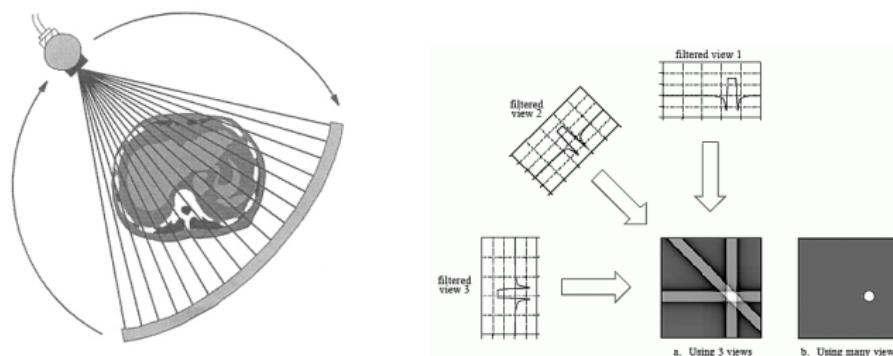
- fdk.py
 - Main file used for benchmarks
- fdkvisual.py
 - Main file used for visualization
- fdkload.py
 - Module used for loading data
- fdkcore.py
 - Core reconstruction kernel to be modified

FDK core kernel

```
1 for p in xrange(nr_projections):
2     for z in xrange(z_voxels):
3         combined_matrix[2, :] = z voxel_coords[z]
4
5         vol_det_map = dot(transform_matrix[p], combined_matrix)
6
7         map_cols = rint(divide(vol_det_map[0, :],
8                             vol_det_map[2, :])).astype(int32)
9
10        map_rows = rint(divide(vol_det_map[1, :],
11                            vol_det_map[2, :])).astype(int32)
12
13        mask = (map_cols >= 0) & (map_rows >= 0) & \
14            (map_cols < detector_columns) & \
15            (map_rows < detector_rows)
16
17        proj_indexs = map_cols * mask + \
18                      map_rows * mask * detector_columns
19
20        recon_volume[z].flat += flat_proj_data[proj_indexs] * \
21                               volume_weight[p] * mask
```

FDK predefined data

- projections.bin
 - Matrix with weighted and filtered 2D X-Ray images



FDK predefined data

- combined.bin
 - Matrix with all combinations of all X,Y coordinates for the 3D

| | Voxel 0 | Voxel 1 | Voxel ... | Voxel X*Y |
|-------|---------|---------|-----------|-----------|
| X | val | val | ... | val |
| Y | val | val | ... | val |
| Z | - | - | - | - |
| Dummy | 1 | 1 | 1 | 1 |

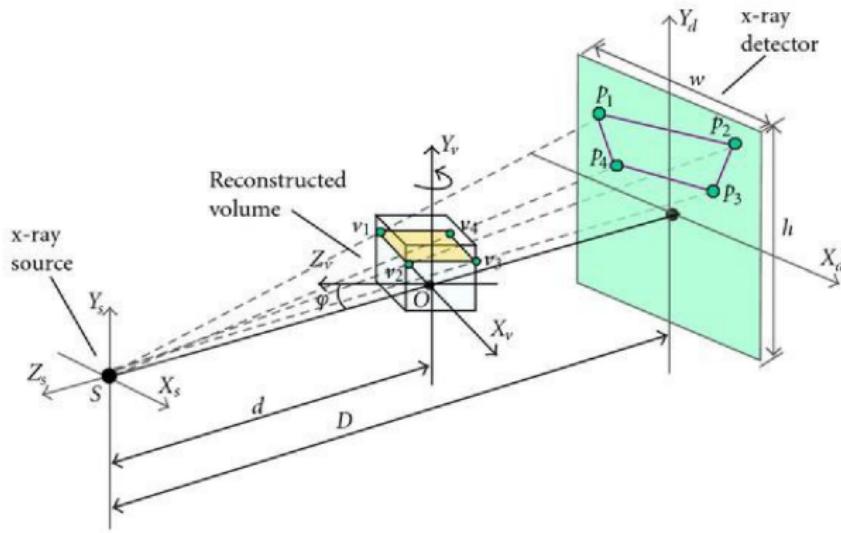
FDK predefined data

- z voxel coords bin
 - Matrix with Z coordinates for the 3D volume

| | Voxel 0 | Voxel 1 | Voxel ... | Voxel X*Y |
|---|---------|---------|-----------|-----------|
| Z | val | val | ... | val |

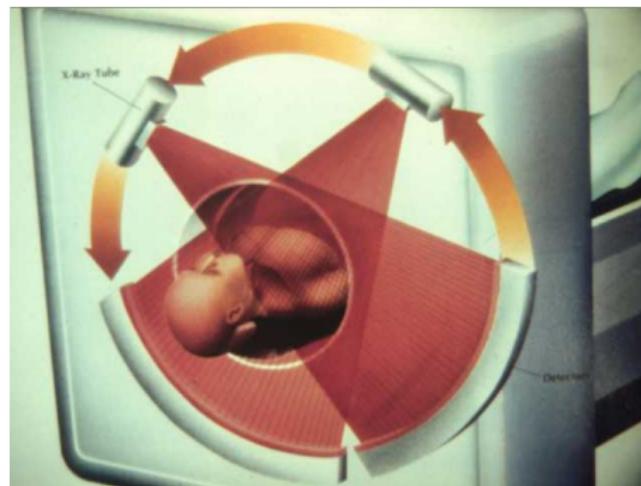
FDK predefined data

- transform.bin
 - Matrix used to map 3D coordinates to 2D coordinates



FDK predefined data

- volumeweight.bin
 - Post weight to compensate for the cone effect of the X-Ray beam

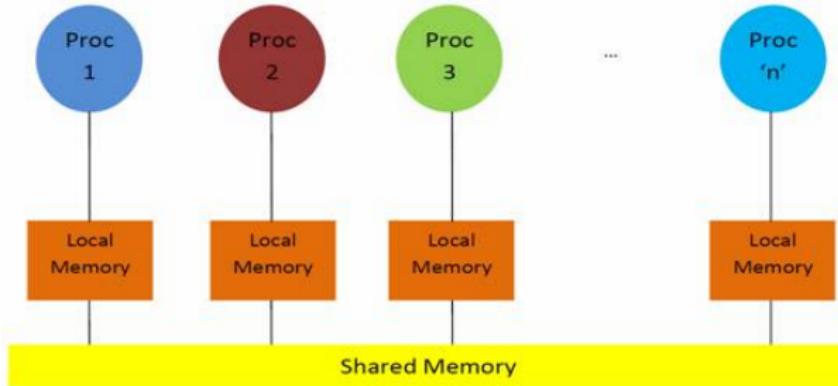


FDK core kernel

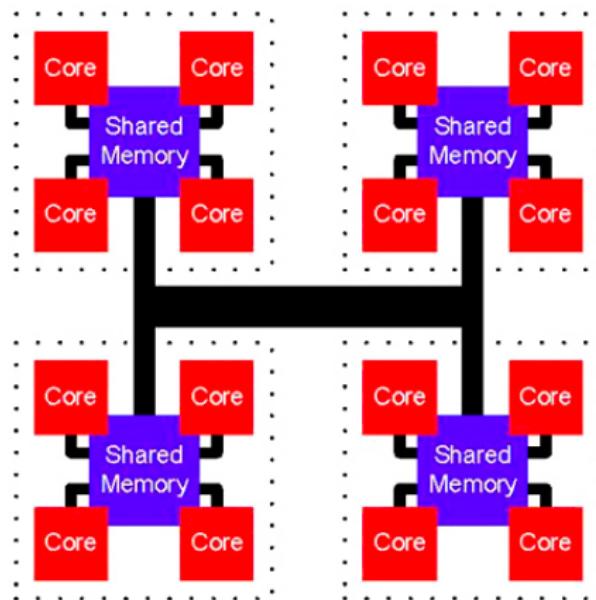
```
1 for p in xrange(nr_projections):
2     for z in xrange(z_voxels):
3         combined_matrix[2, :] = z voxel_coords[z]
4
5         vol_det_map = dot(transform_matrix[p], combined_matrix)
6
7         map_cols = rint(divide(vol_det_map[0, :],
8                             vol_det_map[2, :])).astype(int32)
9
10        map_rows = rint(divide(vol_det_map[1, :],
11                            vol_det_map[2, :])).astype(int32)
12
13        mask = (map_cols >= 0) & (map_rows >= 0) & \
14            (map_cols < detector_columns) & \
15            (map_rows < detector_rows)
16
17        proj_indexs = map_cols * mask + \
18                      map_rows * mask * detector_columns
19
20        recon_volume[z].flat += flat_proj_data[proj_indexs] * \
21                               volume_weight[p] * mask
```

Programming SMP systems

Architecture overview



Architecture overview



Python Multiprocessing

```
1 from multiprocessing import Process
2 import numpy
3 import shmemarray
4
5 def worker(id, result):
6     '''thread worker function'''
7     result[id, :] = numpy.random.random(10)
8
9 def main():
10    workers = 4
11
12    data = shmemarray.create((workers, 10), numpy.float)
13
14    processes = [Process(target=worker, args=(i, data)) \
15                 for i in xrange(workers)]
16
17    for p in processes: p.start()
18    for p in processes: p.join()
19
20 if __name__ == '__main__':
21    main()
```

Python Multiprocessing

```
1 from multiprocessing import Process, Lock
2 import numpy
3 import shmemarray
4
5 def worker(id, result, lock):
6     '''thread worker function'''
7     lock.acquire(True)
8     result[:] += numpy.random.random(10)
9     lock.release()
10
11 def main():
12     lock = Lock()
13     workers = 4
14
15     data = shmemarray.create((10), numpy.float)
16
17     processes = [Process(target=worker, args=(i, data, lock)) \
18                  for i in xrange(workers)]
19
20     for p in processes: p.start()
21     for p in processes: p.join()
22
23 if __name__ == '__main__':
24     main()
```

Python Multiprocessing

```
1 def worker(id, result):
2     '''thread worker function'''
3     result[id, :] = numpy.random.random(10)
```

```
1 def worker(id, result, lock):
2     '''thread worker function'''
3     lock.acquire(True)
4     result[:] += numpy.random.random(10)
5     lock.release()
```

The Assignment

- SMP: Simpler programming model than OpenCL
- But expect a lot smaller speedup

Recon Execution time 256x256x256

Intel(R) Core(TM) i7-3770 CPU @ 3.40GHz: 488.8565s

NVIDIA GTX-680: 1.8033s

NVIDIA GTX-680: 1.0250s

The Assignment

- SMP: Simpler programming model than OpenCL
- But expect a lot smaller speedup

Recon Execution time 256x256x256

Intel(R) Core(TM) i7-3770 CPU @ 3.40GHz: 488.8565s

NVIDIA GTX-680: 1.8033s

NVIDIA GTX-680: 1.0250s

The Assignment

- SMP: Simpler programming model than OpenCL
- But expect a lot smaller speedup

Recon Execution time 256x256x256

Intel(R) Core(TM) i7-3770 CPU @ 3.40GHz: 488.8565s

NVIDIA GTX-680: 1.8033s

NVIDIA GTX-680: 1.0250s

The Assignment

- SMP: Simpler programming model than OpenCL
- But expect a lot smaller speedup

Recon Execution time 256x256x256

Intel(R) Core(TM) i7-3770 CPU @ 3.40GHz: 488.8565s

NVIDIA GTX-680: 1.8033s

NVIDIA GTX-680: 1.0250s

The Assignment

- SMP: Simpler programming model than OpenCL
- But expect a lot smaller speedup

Recon Execution time 256x256x256

Intel(R) Core(TM) i7-3770 CPU @ 3.40GHz: 488.8565s

NVIDIA GTX-680: 1.8033s

NVIDIA GTX-680: 1.0250s

The Assignment

- SMP: Simpler programming model than OpenCL
- But expect a lot smaller speedup

Recon Execution time 256x256x256

Intel(R) Core(TM) i7-3770 CPU @ 3.40GHz: 488.8565s

NVIDIA GTX-680: 1.8033s

NVIDIA GTX-680: 1.0250s

The Assignment

To be released the 25th February at 15.00

Remember:

This is a SMP assignment

The Assignment

To be released the 25th February at 15.00

Remember:

This is a SMP assignment

The Assignment

To be released the 25th February at 15.00

Remember:

This is a SMP assignment